

Priority Queues (Heaps) Data Structure


Queues : Limitations

- Consider applications
 - ordering CPU jobs
 - searching for the exit in a maze
 - emergency room admission processing
- Problems?
 - short jobs **should go first**
 - most promising nodes **should be searched first**
 - most urgent cases **should go first**

Priority Queue ADT

- Priority Queue operations

- create
- destroy
- insert
- deleteMin
- is_empty



- Priority Queue property: for two elements in the queue, x and y , if x has a lower **priority value** than y , x will be deleted before y

Applications of the Priority Q

- Hold jobs for a printer in order of length
- Store packets on network routers in order of urgency
- Simulate events
- Anything *greedy*

Naïve Priority Queue Data Structures

- Linked list:
 - *Insert: $O(1)$*
 - *DeleteMin: $O(N)$*
- Sorted Linked list:
 - *Insert: $O(N)$*
 - *DeleteMin: $O(1)$*

BST Tree Priority Queue Data Structure

- Regular BST:

- Insert: $O(\log n)$*


- deleteMin: $O(\log n)$*

- (Average)*

- AVL Tree:


- insert: $O(\log n)$*

- deleteMin: $O(\log n)$*




Binary Heap Priority Q Data Structure


- Heap-order property
 - parent's key is less than children's keys
 - result: minimum is always at the top
- Structure property
 - Binary tree that is completely filled, with the possible exception of the bottom level, which is filled from left to right. Such a tree is known as Complete Binary tree.
 - Complete binary tree of height 'h' has between 2^h and $2^{(h+1)} - 1$ nodes. Height 'h' is $O(\log n)$.



Complete Binary Tree Example




Binary heap Example




Satisfies both structure property and heap property

Basic Operations :Insert

`pqueue.insert(3)`




Percolate Up




DeleteMin

`pqueue.deleteMin()`



Percolate Down




Performance of Binary Heap

	Binary heap worst case	Binary heap avg case	AVL tree worst case	BST tree avg case
Insert	$O(\log n)$	$O(1)$ percolates 1.6 levels	$O(\log n)$	$O(\log n)$
Delete Min	$O(\log n)$	$O(\log n)$	$O(\log n)$	$O(\log n)$

- In practice: binary heaps much simpler to code, lower constant factor overhead

d-Heaps

- Generalization of Binary heap
- Each node has d children
- Still representable by array
- Good choices for d :
 - optimize performance based on # of inserts/removes
 - choose a power of two for efficiency
 - fit one set of children in a cache line
 - fit one set of children on a memory page/disk block
- Insert – $O(\log_d n)$ and Delete $O(n \log_d n)$



New Operation: Merge

Merge(H1,H2): Merge two heaps H1 and H2 of size $O(N)$.

- *E.g.* Combine queues from two different sources to run on one CPU.
- Merging is harder in Binary Heap
- To support merge, we will discuss other data structures

Leftist Heaps

- An alternative heap structure that also enables fast merges
- Based on binary trees rather than k -ary trees


Leftist Heaps

- Leftist heap:
 - almost all nodes are on the left
 - all the merging work is on the right

Null Path Length

the *null path length* (*npl*) of a node is the number of nodes between it and a null in the tree

- $\text{npl}(\text{null}) = -1$
- $\text{npl}(\text{leaf}) = 0$
- $\text{npl}(\text{single-child node}) = 0$




Leftist Heap Properties


- Heap-order property
 - parent's priority value is \leq to childrens' priority values
 - result: minimum element is at the root
- Leftist property
 - null path length of left subtree is \geq npl of right subtree
 - result: tree is at least as “heavy” on the left as the right

Leftist tree examples


NOT leftist



leftist



leftist




every subtree of a leftist
tree is leftist

Skew Heaps

- Problems with leftist heaps
 - extra storage for npl
 - two pass merge (with stack!)
 - extra complexity/logic to maintain and check npl
- Solution: skew heaps
 - blind adjusting version of leftist heaps
 - **amortized** time for merge, insert, and deleteMin is $O(\log n)$
 - worst case time for all three is $O(n)$
 - merge *always* switches children when fixing right path
 - iterative method has only one pass

Merging Two Skew Heaps



Example

